

Figure 3 · Harmful Compliance Rate by Attack Family
(lower = safer · $n = 200$ per family · dual-classifier ensemble)

